



FANSHAWE

INFO-6003

O/S & Application Security

Week 12



Agenda

- Exam in Next Week
- More Access and File System Control
- More detail on PAM
- Patching
- Logging
- init, Upstart init, systemd init
- Managing Services
- chroot

Next Week

- Final Exam
- FOL Respondus Browser
- 90 minutes long
- 3 long answer questions, also MC and TF

More Access Control

Access Control

- Control and Manage the following:
 - User accounts and system administrator functions
 - Access to files, utilities and services
- SUID/SGID and sudo command allow users to assume root permissions for specific commands
 - Only when needed
- Visudo can be used to manage who can use sudo

Visudo

- visudo is the command used to edit the sudoers file
- Contains a list of users who can perform operations with the sudo command
- You can limit the commands a user can use with sudo (**sudo tcpdump** for example)
- We will be doing this in the lab today

Changing Passwords

- The passwd command can be used to set password options on accounts
 - Length of password
 - Lock or Unlock account
 - Number of days a password can be used
 - Days before password can be changed
 - Expiration date
 - File remain intact, but account is locked
- Similar options to what we saw in Windows

Passwd Options

-a, --all	report password status on all accounts
-d, --delete	delete the password for the named account
-e, --expire	force expire the password for the named account
-i, --inactive INACTIVE	set password inactive after expiration to INACTIVE
-l, --lock	lock the password of the named account
-n, --mindays MIN_DAYS	set minimum number of days before password change to MIN_DAYS
-S, --status	report password status on the named account
-u, --unlock	unlock the password of the named account
-w, --warndays WARN_DAYS	set expiration warning days to WARN_DAYS
-x, --maxdays MAX_DAYS	set maximum number of days before password change to MAX_DAYS

File Security & Integrity

- File Integrity Checkers
 - Monitor system files
 - Creates snapshot of files: a hashed signature (message digest) for each file
 - After an attack, compare post-hack signature with snapshot
 - This allows systems administrator to determine which files were changed
- Tripwire is file integrity checker that can be used for some Linux systems
 - Red Hat and SUSE

Tripwire & AIDE

■ Tripwire

- Tripwire detects changes to file system objects
- Tripwire creates a cryptographic hash of a file and stores the hash on each file scanned in a database
- Hashes created on next scan are compared to stored hash
- Changes could indicate a hacker altering files

■ AIDE

- Advanced Intrusion Detection Environment
- Developed as a free replacement to Tripwire

Samhain

- Samhain
 - Host intrusion detection & integrity checker
 - Host & network scan options
 - uses cryptographic checksums of files to detect modifications,
 - can find rogue SUID executable files anywhere on disk
 - Supports logging to a central server

<http://www.la-samhna.de/samhain/>

More on Pluggable Authentication Modules

PAM

- PAM provides a library containing functions for proper authentication procedures
- Allows for a separate module to provide authentication so it does not have to be in the program API

PAM

- PAM configuration files have 3 entries
 - The first entry indicates one of four categories which identify different types of modules for controlling access to a particular service
 - The second field in each entry is called the control flag and determines the action taken when the module succeeds or fails
 - The third field contains the values

```
password [success=2 default=ignore] pam_unix.so obscure sha512
```

PAM

- Authentication
 - Provides the actual authentication (perhaps asking for and checking a password) and sets credentials, such as group membership or Kerberos tickets
- Account
 - Checks the account has not expired, the user is allowed to log in at this time of day, and so on
- Password
 - Module is used to set passwords
- Session
 - Used after a user has been authenticated
 - Performs additional tasks which are needed to allow access such as mounting the user's home directory

PAM Control Options

- Required
 - The module must succeed
 - Regardless of whether the module fails or succeeds, processing will continue with the next line (other modules of the same module type will be executed)
 - At the end of all of the processing, a failure will be recorded
- Requisite
 - The module must succeed for the module type to succeed
 - If it fails, processing stops immediately
 - If it succeeds, processing continues with the next line

PAM Control Options

- Sufficient
 - If the module succeeds, then the module type succeeds and processing stops immediately
 - If it fails, processing continues with the next line
- Optional
 - The module is executed, but the failure or success of the module is ignored
- Include
 - In place of a module name, another configuration file is given
 - All of the lines of the same type from that configuration file are treated as if they were present in this configuration file

Logging

Logging

- Logging is used to record system events
 - This is sometimes called auditing
- Linux syslogd, rsyslogd and kernel log daemon, klogd, write events to log files
- The `/etc/syslog.conf` or `/etc/rsyslog.conf` file specifies where log files are located
 - Exact location depends on the version of Linux
 - In addition to log location the `.conf` files also specify additional parameters such as log file format

Logging

- `/var/log/secure`
 - Successful & failed logins
- `/var/log/messages`
 - General error messages from kernel or other services
- `/var/log/boot.log`
 - Information logged during system boot
- Specific services that have a lot of activity can have their own log file
 - `/var/log/httpd`
 - Web service messages

Logging

- /var/log or /var/adm hold the following files that are used to investigate hacking or suspicious activity
 - utmp
 - Current status of system: boot time, logins, logouts, system events
 - wtmp
 - Historical record of utmp
 - btmp
 - Records failed login attempts
 - lastlog
 - Last time and location of user's last login to system
 - Console, ssh or telnet

Patching

Update Manager

- Most Linux distributions have an update tool to download patches for OS or software packages
 - Some have more than one package manager
- Red Hat, Fedora, CentOS
 - **up2date** (Red Hat Update Agent) & **YUM** (Yellowdog Updater, Modified), **RPM** (Red Hat Package Mgr.)
- SUSE Linux
 - **YaST** (Yet another Setup Tool)
- Debian
 - **APT** (Advanced Packaging Tool)

Apt Command Examples

- `apt-get update`
 - Updates packages listings from the repo, should be run at least once a week
- `apt-get upgrade`
 - Upgrades all currently installed packages with those updates available from the repo. should be run once a week
- You may need to turn off your AV to get these commands to work
 - Especially when we start downloading hacking tools

Apt Command Examples

- `apt-cache search <pattern>`
 - Searches packages and descriptions for <pattern>
- `apt-get install <package>`
 - Downloads <package> and all of its dependencies, and installs or upgrades them
- `apt-cache search extundelete`
 - Will search the repository for extundelete
- `apt-cache install extundelete`
 - Will install extundelete on your system

Hardening or Hardened

Hardening or Hardened

- When it comes to securing your Linux distribution you have a couple choices
- Hardening
 - This is the process of locking down an existing distribution to make it more secure
- Hardened
 - Choosing a distribution that has been specifically developed with security in mind

Hardening or Hardened

- EnGarde Secure Linux (hardened distro)
 - <http://www.engardelinux.org>
- Hardened Linux (hardened distro)
 - <http://hardenedlinux.sourceforge.net>
- Bastille (hardening program)
 - <http://bastille-linux.sourceforge.net>
- SELinux (security enhancements)
 - <http://selinuxproject.org>
- AppArmor (security enhancements)
 - <http://wiki.apparmor.net>

Bastille

- The Bastille hardening program
 - Designed to lock a system down based on best practices
 - Interactively configures the system
 - Can be used to analyze the systems state
 - Supports a wide variety of platforms
 - Red Hat / Fedora Core / SUSE / Mandrake
 - Debian
 - Gentoo
 - HP-UX
 - Mac OS X

SELinux

- SELinux is not a hardened distribution
 - It is a set of Kernel modifications and tools that can be added to Linux distributions
- Primarily developed by the NSA
- Makes heavy use of MACs and Linux Security Modules (LSMs) in the Linux Kernel
 - The MACs are used in place of DACs
- Great for hardening the system, but does make management more difficult

SELinux

- Makes use of subjects, objects, labels and policies to lock the system down
- Subjects
 - Users, applications, process, etc.
- Objects
 - Files, sockets
- Labels
 - Metadata applied to objects
- Policies
 - Access permissions for subjects and objects

Init Daemons

init Daemon

- The first process to start, aside from the kernel, is the init daemon
 - The Linux Kernel has a process ID of 0
 - The init daemon has a process ID of 1
- The init daemon starts all the other processes and services as the system boots
- The traditional init isn't used in modern distros
 - System V init
- Modern versions of init
 - Upstart init
 - systemd init

init Daemon

- The newer init Daemons were developed to deal with dynamic environments
 - USB keys
 - Other hot-plug devices
- The newer init Daemons have some backwards compatibility built in, but the move is towards getting rid of this

Run Levels

- Allow users to start the OS in different operating modes
 - Depending on the run level different functionality will be available
 - Single User
 - Multi-User
 - Network Services start, or Stopped
- Windows also has this functionality, but implemented differently
 - Safe Mode, Safe Mode with Networking, Command Prompt, etc.

Run Levels

- Run Level 0
 - Used to Halt or shutdown the system
- Run Level 1
 - Single user mode minimal configuration
- Run Level 2
 - Multi user mode without networking
- Run Level 3
 - Multi user mode with networking
- Run Level 5
 - X Window mode (Full GUI)
- Run Level 6
 - Used to reboot the system D

Variations in Run Levels

- Run levels vary by distribution
 - Especially run levels 2 through 4
- You can use the `init` command to change run levels (good for troubleshooting)
 - `init 0`: Will shutdown the system
 - `init 1`: Will start the system in single user mode
 - `init 6`: Will reboot the system
- You can see your current run level a couple ways
 - `runlevel`
 - `who -r`

init

init

- Uses the `/etc/inittab` file to determine the default run level
 - This is used if you don't actually specify a run level
 - `id:5:initdefault:`
 - Would start an X windows sessions (GUI)
- Based on the run level, certain scripts are run
 - `/etc/rc.d/rc#.d`
 - `/etc/rc.d/rc5.d` in the example above
 - The location of the files varies by distro

init rc#.d

- The rc#.d directories contain scripts that start or kill processes
 - S for starting process
 - K for killing process
- The scripts are also numbered which allows them to be started or stopped in a particular order

```
K73winbind  
K73ypbind  
K74nscd  
K74ntpd  
K84btseed  
K84bttrack  
K87multipathd
```

```
S13mcstrans  
S13rpcbind  
S13setroubleshoot  
S14nfslock  
S15mdmonitor  
S18rpcidmapd  
S19rpcgssd
```

```
S97dhcdbd  
S97yum-updatesd  
S98avahi-daemon  
S98haldaemon  
S99firstboot  
S99local  
S99smartd
```


Dependencies

- The fact that the administrator can control the order the services start allows dependencies to be taken into account
 - If service A is a dependency for service B it will need to have a lower number, so that service A is up and running before service B

Upstart init

Upstart init

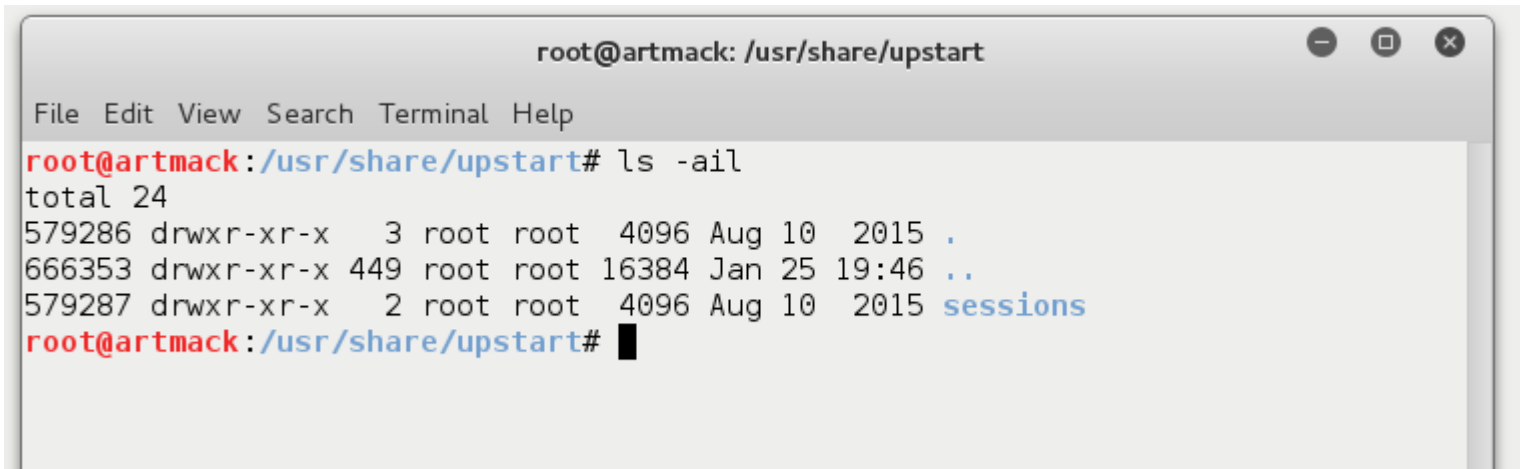
- Used in many newer distributions
 - Fedora 9 to 14
 - Ubuntu 6.10 and later
 - Google's Chrome OS
- The primary difference between Upstart init and init is the handling of starting and stopping services
 - Geared towards handling and ever changing environment
 - Hot pluggable devices

Upstart init

- Handles services through defined jobs
 - Jobs can be either a task of a service
- Tasks
 - Performs a limited duty, and when finished returns to the waiting state
 - stop/waiting
- Service
 - Long running program that never self-terminates
 - Stays in the running state
 - start/running

Upstart init

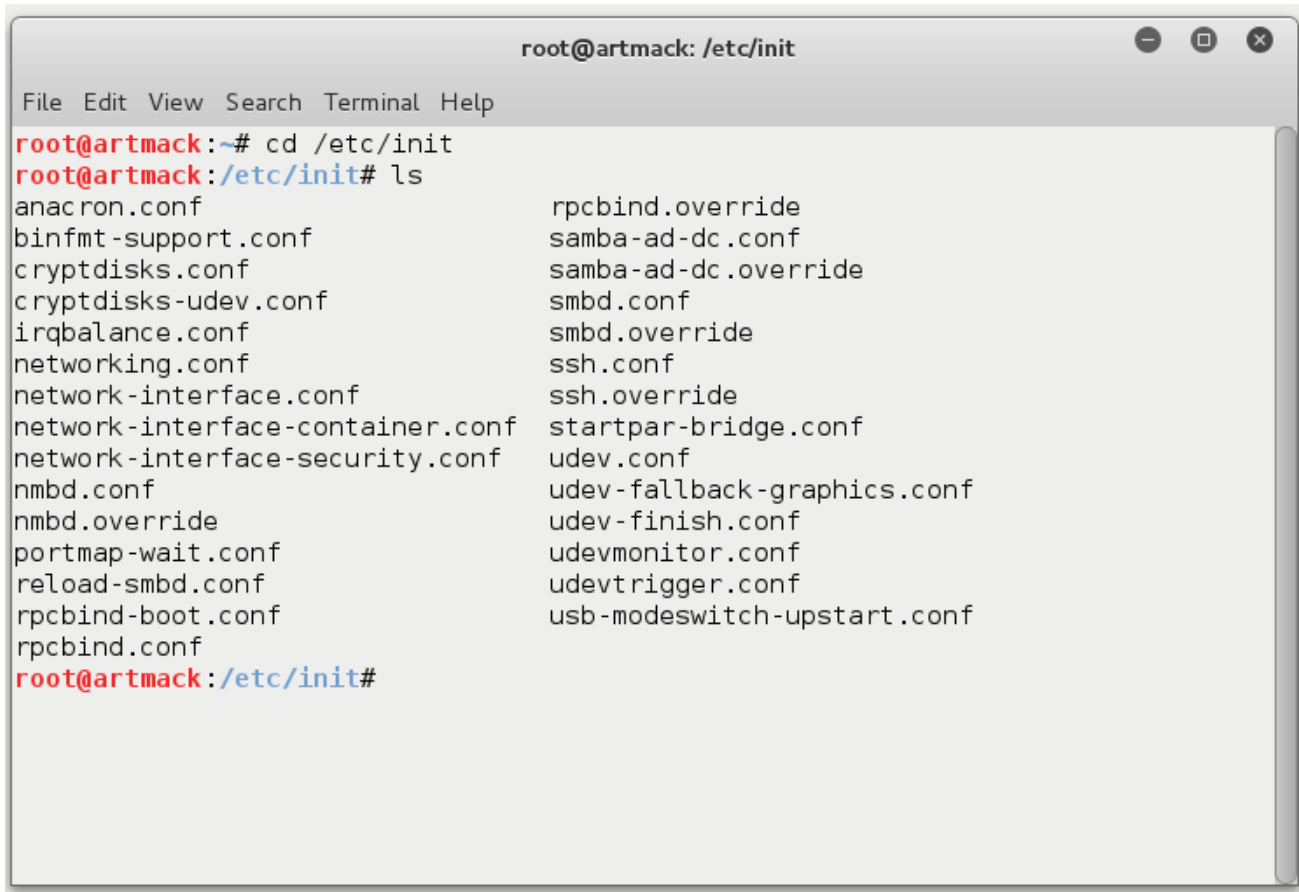
- /usr/share/upstart
- Good way of finding out that you are on an Upstart-based system



```
root@artmack: /usr/share/upstart
File Edit View Search Terminal Help
root@artmack:/usr/share/upstart# ls -ail
total 24
579286 drwxr-xr-x  3 root root  4096 Aug 10  2015 .
666353 drwxr-xr-x 449 root root 16384 Jan 25 19:46 ..
579287 drwxr-xr-x  2 root root  4096 Aug 10  2015 sessions
root@artmack:/usr/share/upstart#
```

Upstart init

- The various jobs are defined in /etc/init



```
root@artmack: /etc/init
File Edit View Search Terminal Help
root@artmack:~# cd /etc/init
root@artmack:/etc/init# ls
anacron.conf                      rpcbind.override
binfmt-support.conf              samba-ad-dc.conf
cryptdisks.conf                  samba-ad-dc.override
cryptdisks-udev.conf             smbd.conf
irqbalance.conf                  smbd.override
networking.conf                  ssh.conf
network-interface.conf           ssh.override
network-interface-container.conf startpar-bridge.conf
network-interface-security.conf  udev.conf
nmbd.conf                        udev-fallback-graphics.conf
nmbd.override                    udev-finish.conf
portmap-wait.conf                udevmonitor.conf
reload-smbd.conf                  udevtrigger.conf
rpcbind-boot.conf                usb-modeswitch-upstart.conf
rpcbind.conf
root@artmack:/etc/init#
```

Upstart init

- The job definition files determine what services are started, stopped, restarted, etc.
- When certain conditions are met certain actions are taken
 - Allows for a much more dynamic environment
 - Ability to respond to change

systemd init

systemd init

- Alternative to init and Upstart init
- Runs on a variety of distributions
 - Arch Linux (since 2012)
 - Fedora 15 and later
 - Mandriva (since 2011)
 - openSUSE 12.1 and later
- Some key differences as compared to Upstart init
 - Starts fewer services
 - Starts services in a parallel manner
 - Supervises all processes

systemd init

- Instead of run levels systemd uses target units
- A unit is a group consisting of
 - name
 - type
 - configuration file
- A unit is focused on a particular service or action

systemd init

- There are eight unit types
 - automount
 - device
 - mount
 - path
 - service
 - snapshot
 - socket
 - target
- The service and target units primarily deal with services

systemd init

- A service unit is used to manage a daemon
 - rsyslog.service
 - sshd.service
- A target unit is a group of other units
 - sysinit.target
 - All the actions required for system initialization
 - syslog.target

Auditing Services

Auditing Services

- For init systems

- `chkconfig –list`

- Lets you know what services are on or off for particular run levels

- `service –status-all`

- Allows you to see if a service is running or not

- For Upstart init systems

- `initctl list`

- Let you see the service state

- start/running

- stop/waiting

Auditing Services

- For systemd init systems
 - `systemctl list-unit-files --type=unittype`
- With all of the auditing commands you can use `grep` to filter your results
 - `| grep running`
 - for running services
 - `| grep wait`
 - for stopped services

Controlling Services

Controlling Services

- All the init daemons have built in mechanisms for manually controlling/polling the services
 - start
 - stop
 - restart (will stop and start the service)
 - reload (only reloads the configuration file)
 - status
- These can be helpful when troubleshooting
- Additionally, you can restart a service without rebooting the server

Controlling Services

- With init
 - service cups status
 - service cups start
 - service cups stop
 - service cups restart
- With Upstart init
 - initctl status cups
 - initctl start cups
 - initctl stop cups
 - initctl restart cups
 - initctl reload cups

Controlling Services

- With systemd init
 - `systemctl status cups.service`
 - `systemctl start cups.service`
 - `systemctl stop cups.service`
 - `systemctl restart cups.service`
 - `systemctl reload cups.service`
 - `systemctl condrestart cups.service`
- The `condrestart` is a conditional restart that will only restart the service if it is already running
 - inactive services stay inactive

chroot

chroot

- Pronounced “cha-root”
- This is also sometimes referred to as “chroot jail”
- Chroot will limit access to only that part of the file system defined by chroot
 - Normally implemented to restrict access for untrusted or anonymous users, untrusted applications etc.
 - The initial process and all its child processes will perceive the root directory to be that which is defined when setting up the chroot environment

chroot

- The directory configured for chroot is treated as the root of the files system for all processes started in the chroot directory
 - Impossible to access any files or binaries outside of the chroot directory
 - As long as no privilege escalation exists
- Because an application running within a chroot jail can't access any files outside, all its required files need to be within the chroot jail
 - passwd files
 - libraries
 - binaries, etc.

chroot

- Directory Structure

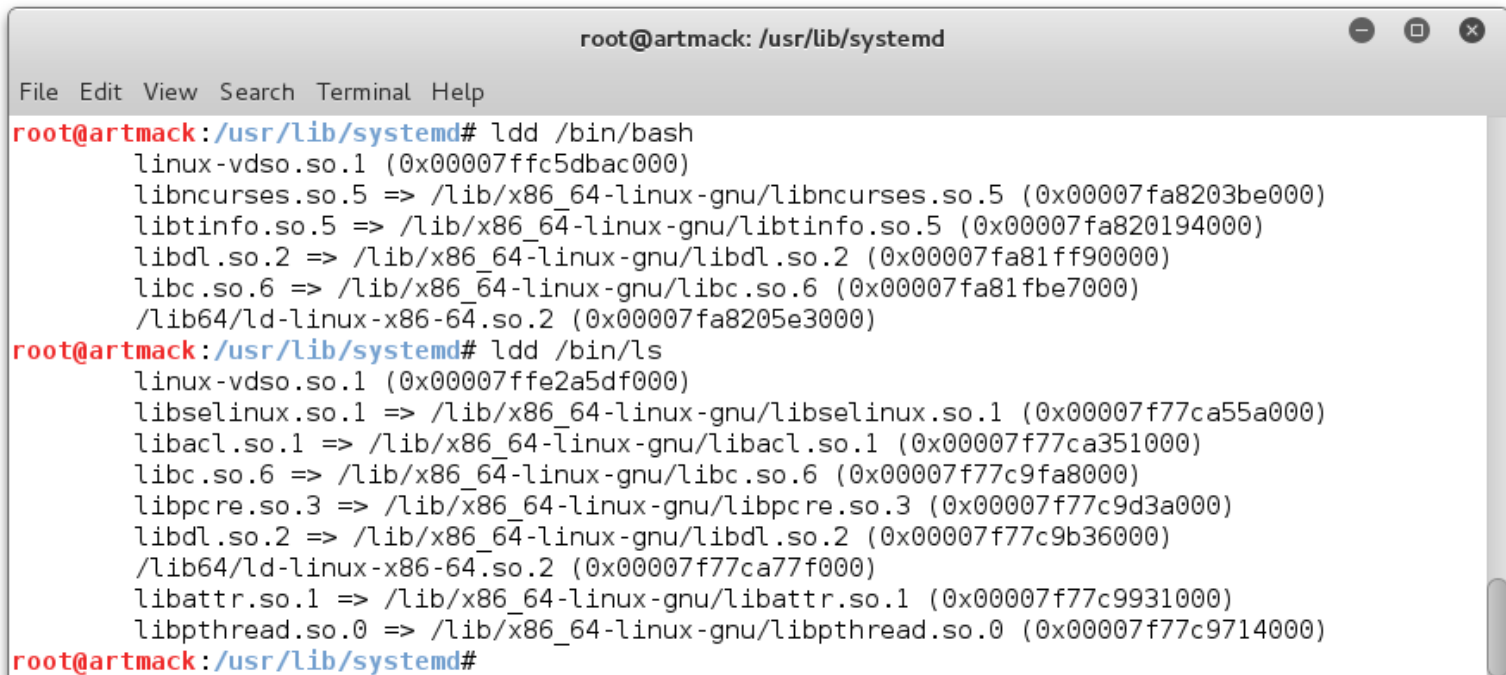
- You match the expected directory structure within the chroot environment

- Bash and ls example

- If you wanted to run a bash shell and use the ls command in a chroot environment you would need the following
 - /tmp/newroot/bin/bash
 - to hold bash and ls
 - /tmp/newroot/lib
 - to hold the libraries

chroot

- You can determine which libraries will be required to run your programs
 - Will vary by distribution



```
root@artmack: /usr/lib/systemd
File Edit View Search Terminal Help
root@artmack:/usr/lib/systemd# ldd /bin/bash
linux-vdso.so.1 (0x00007ffc5dbac000)
libncurses.so.5 => /lib/x86_64-linux-gnu/libncurses.so.5 (0x00007fa8203be000)
libtinfo.so.5 => /lib/x86_64-linux-gnu/libtinfo.so.5 (0x00007fa820194000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007fa81ff90000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fa81fbe7000)
/lib64/ld-linux-x86-64.so.2 (0x00007fa8205e3000)
root@artmack:/usr/lib/systemd# ldd /bin/ls
linux-vdso.so.1 (0x00007ffe2a5df000)
libselinux.so.1 => /lib/x86_64-linux-gnu/libselinux.so.1 (0x00007f77ca55a000)
libacl.so.1 => /lib/x86_64-linux-gnu/libacl.so.1 (0x00007f77ca351000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f77c9fa8000)
libpcre.so.3 => /lib/x86_64-linux-gnu/libpcre.so.3 (0x00007f77c9d3a000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007f77c9b36000)
/lib64/ld-linux-x86-64.so.2 (0x00007f77ca77f000)
libattr.so.1 => /lib/x86_64-linux-gnu/libattr.so.1 (0x00007f77c9931000)
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00007f77c9714000)
root@artmack:/usr/lib/systemd#
```


chroot

- After copying all the bin, ls and library files to the chroot environment you need to enable to environment
- Most distributions contain a program called chroot that invokes the chroot() system call for you
 - The program takes two variable
 - chroot directory
 - command we want to run
- In our example
`chroot /tmp/newroot /bin/bash`

chroot


- It can be difficult to find all the libraries and other dependencies required
- There are tools available to find these dependencies
 - ldd finds the library dependencies
 - strace finds the system calls
 - lsof lists the open files and the processes that opened them

chroot Cautions

- Never add any files into chroot that have functions that can be used to escape
 - No compilers
 - No interpreter
 - No services that require root to run
- Daemons placed in chroot should not run with root permissions
 - They would be able to break out

Homework

- Read Online Chapter 25
- 25.1 – 25.9 – Linux Security



Lab 10 – initCtl & Sudo

Lab 10 Details

- Create LAN Segment for VMs
- Configure Network Interfaces
- Use Sudo for administration